ECE M146 HW6
Melody Chen
705120273

1. The pdf for two jointly Gaussian random variables X and Y is of the following form parameterized by the scalars m_1, m_2, σ_1, σ_2 and ρ_{XY}:

$$f_{X,Y}(x,y) = \frac{\exp\left\{\frac{-1}{2(1-\rho_{XY}^2)}\left[\left(\frac{x-m_1}{\sigma_1}\right)^2 - 2\rho_{XY}\left(\frac{x-m_1}{\sigma_1}\right)\left(\frac{y-m_2}{\sigma_2}\right) + \left(\frac{y-m_2}{\sigma_2}\right)^2\right]\right\}}{2\pi\sigma_1\sigma_2\sqrt{1-\rho_{XY}^2}}. \quad (1)$$

The pdf for multivariate jointly Gaussian random variable $Z \in \mathbb{R}^k$ is of the following form parameterized by $\mu \in \mathbb{R}^k$ and $\Sigma \in \mathbb{R}^{k\times k}$.

$$f_Z(z) = \frac{\exp\left\{-\frac{1}{2}(z-\mu)^T\Sigma^{-1}(z-\mu)\right\}}{\sqrt{(2\pi)^k|\Sigma|}}. \quad (2)$$

Suppose $Z = [X,Y]^T$, i.e., $z = [x,y]^T$.

(a) Find μ, Σ^{-1} and Σ in terms of m_1, m_2, σ_1, σ_2 and ρ_{XY}.

(b) Suppose $\rho_{XY} = 0$, what is Σ in this case? Can you write $f_{X,Y}(x,y)$ as the product of two single variate Gaussian distributions? Are X and Y independent?

a) By comparing the two equations above, we see that when multivariate jointly Gaussian Density is applied to two variables:

$$\mu = \begin{bmatrix} m_1 \\ m_2 \end{bmatrix}$$

$$\Sigma^{-1} = \frac{1}{\sigma_1^2\sigma_2^2(1-\rho_{XY})^2}\begin{bmatrix} \sigma_2^2 & -\rho_{XY}\sigma_1\sigma_2 \\ -\rho_{XY}\sigma_1\sigma_2 & \sigma_1^2 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} \sigma_1^2 & \rho_{XY}\sigma_1\sigma_2 \\ \rho_{XY}\sigma_1\sigma_2 & \sigma_2^2 \end{bmatrix}.$$

We double check by plugging μ, Σ^{-1}, Σ in to equation (2):

$$f_Z\left(\begin{bmatrix} x \\ y \end{bmatrix}\right) = \frac{\exp\left\{-\frac{1}{2}\left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} m_1 \\ m_2 \end{bmatrix}\right)^T \cdot \frac{1}{\sigma_1^2\sigma_2^2(1-\rho_{XY})^2}\begin{bmatrix} \sigma_2^2 & -\rho_{XY}\sigma_1\sigma_2 \\ -\rho_{XY}\sigma_1\sigma_2 & \sigma_1^2 \end{bmatrix} \cdot \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} m_1 \\ m_2 \end{bmatrix}\right)\right\}}{\sqrt{(2\pi)^2(\sigma_1^2\sigma_2^2 - \rho_{XY}^2\sigma_1^2\sigma_2^2)}}$$

$$= \frac{\exp\left\{-\frac{1}{2}\frac{1}{\sigma_1^2\sigma_2^2(1-\rho_{XY})^2}\left[(x-m_1)\sigma_2^2 - (y-m_2)\rho_{XY}\sigma_1\sigma_2, -(x-m_1)\rho_{XY}\sigma_1\sigma_2 + (y-m_2)\sigma_1^2\right]\begin{bmatrix} x-m_1 \\ y-m_2 \end{bmatrix}\right\}}{2\pi\sigma_1\sigma_2\sqrt{(1-\rho_{XY}^2)}}$$

$$= \frac{\exp\left\{-\frac{1}{\sigma_1^2\sigma_2^2(1-\rho_{XY})^2}\left[(x-m_1)^2\sigma_2^2 - (y-m_2)(x-m_1)\rho_{XY}\sigma_1\sigma_2 + (-(x-m_1)(y-m_2)\rho_{XY}\sigma_1\sigma_2 + (y-m_2)^2\sigma_1^2)\right]\right\}}{2\pi\sigma_1\sigma_2\sqrt{(1-\rho_{XY}^2)}}$$

$$= \frac{\exp\left\{-\frac{1}{(1-\rho_{XY})^2}\left[\left(\frac{x-m_1}{\sigma_1}\right)^2 - 2\rho_{XY}\frac{(x-m_1)(y-m_2)}{\sigma_1\sigma_2} + \left(\frac{y-m_2}{\sigma_2}\right)^2\right]\right\}}{2\pi\sigma_1\sigma_2\sqrt{(1-\rho_{XY}^2)}}$$

b) If $\rho_{XY}=0$, $\Sigma = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}$ $\quad f_{X,Y}(x,y) = \frac{\exp\left\{-\frac{1}{2}\left[\left(\frac{x-m_1}{\sigma_1}\right)^2 + \left(\frac{y-m_2}{\sigma_2}\right)^2\right]\right\}}{2\pi\sigma_1\sigma_2}$

$$= \frac{1}{2\pi\sigma_1\sigma_2} e^{-\frac{1}{2}\left(\frac{x-m_1}{\sigma_1}\right)^2} \cdot e^{-\frac{1}{2}\left(\frac{y-m_2}{\sigma_2}\right)^2}$$

$$= f(x)\cdot f(y)$$

Yes, X and Y are independent as $f(x,y) = f(x)\cdot f(y)$.

2. The Gaussian Discriminant Analysis (GDA) models the class conditional distribution as multivariate Gaussian, i.e, $P(X|Y) \sim \mathcal{N}(\mu_Y, \Sigma)$. Suppose we want to enforce the **Naive Bayes (NB) assumption**, i.e. $P(X_i|Y, X_j) = P(X_i|Y), \forall j \neq i$, to GDA. Show that all off diagonal elements of Σ equal to 0: $\Sigma_{i,j} = 0, \forall i \neq j$ with the **NB assumption**.

$Cov(x,y) = E((x - E(x))(y - E(y)))$ by definition.

$Cov(x_i|y, x_j|y) = E((x_i|y - E(x_i|y))(x_j|y - E(x_j|y))$

$= E((x_i|y)(x_j|y) - (x_i|y)E(x_j|y) - (x_j|y)E(x_i|y) + E(x_i|y)E(x_j|y))$

$= 2E(x_i|y)E(x_j|y) - 2E(x_i|y)E(x_j|y)$

$= 0$

* $E(E(x_j|y)) = E(x_j|y)$

3. Consider the classification problem for two classes, C_0 and C_1. In the generative approach, we model the class-conditional distribution $P(x|C_0)$ and $P(x|C_1)$, as well as the class priors $P(C_0)$ and $P(C_1)$.The posterior probability for class C_0 can be written as

$$P(C_0|x) = \frac{P(x|C_0)P(C_0)}{P(x|C_0)P(C_0) + P(x|C_1)P(C_1)}.$$

(a) Show that $P(C_0|x) = \sigma(a)$ where $\sigma(a)$ is the *sigmoid* function defined by

$$\sigma(a) = \frac{1}{1 + \exp(-a)}.$$

Find a in terms of $P(x|C_0)$, $P(x|C_1)$, $P(C_0)$ and $P(C_1)$.

$$P(C_0|x) = \frac{1}{1 + \frac{P(x|C_1)P(C_1)}{P(x|C_0)P(C_0)}}$$

$$\exp(-a) = \frac{P(x|C_1)P(C_1)}{P(x|C_0)P(C_0)}$$

$$a = -\ln\left(\frac{P(x|C_1)P(C_1)}{P(x|C_0)P(C_0)}\right) = \ln\left(\frac{P(x|C_0)P(C_0)}{P(x|C_1)P(C_1)}\right)$$

(b) In the GDA model, we have the class conditional distribution as follows

$$P(x|C_0) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0)\right),$$

$$P(x|C_1) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1)\right).$$

Suppose we are able to find the maximum likelihood estimation of $\mu_0, \mu_1, \Sigma, P(C_0)$, and $P(C_1)$. Show that $a = w^Tx + b$ for some w and b. Find w and b in terms of $\mu_0, \mu_1, \Sigma, P(C_0)$, and $P(C_1)$. This shows that the decision boundary is linear.

$$a = \ln\left(\frac{\frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \cdot \exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0)\right) \cdot P(C_0)}{\frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \cdot \exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1)\right) \cdot P(C_1)}\right)$$

$$= \ln\left(\exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0)\right) \cdot P(C_0)\right) - \ln\left(\exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1)\right) \cdot P(C_1)\right)$$

$$= \left(-\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0)\right) + \ln(P(C_0)) - \left(-\frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1)\right) + \ln(P(C_1))$$

$$= -\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0) + \frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= -\frac{1}{2}(x^T-\mu_0^T)\Sigma^{-1}(x-\mu_0) + \frac{1}{2}(x^T-\mu_1^T)\Sigma^{-1}(x-\mu_1) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= -\frac{1}{2}(x^T\Sigma^{-1} - \mu_0^T\Sigma^{-1})\cdot(x-\mu_0) + \frac{1}{2}(x^T\Sigma^{-1} - \mu_1^T\Sigma^{-1})(x-\mu_1) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= -\frac{1}{2}(x^T\Sigma^{-1}x - x^T\Sigma^{-1}\mu_0 - \mu_0^T\Sigma^{-1}x + \mu_0^T\Sigma^{-1}\mu_0) + \frac{1}{2}(x^T\Sigma^{-1}x - x^T\Sigma^{-1}\mu_1 - \mu_1^T\Sigma^{-1}x + \mu_1^T\Sigma^{-1}\mu_1) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= -\frac{1}{2}(-2x^T\Sigma^{-1}\mu_0 + \mu_0^T\Sigma^{-1}\mu_0) + \frac{1}{2}(-2x^T\Sigma^{-1}\mu_1 + \mu_1^T\Sigma^{-1}\mu_1) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= x^T\Sigma^{-1}\mu_0 - \frac{1}{2}\mu_0^T\Sigma^{-1}\mu_0 - x^T\Sigma^{-1}\mu_1 + \frac{1}{2}\mu_1^T\Sigma^{-1}\mu_1 + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= x^T\Sigma^{-1}(\mu_0-\mu_1) \underbrace{- \frac{1}{2}\mu_0^T\Sigma^{-1}\mu_0 + \frac{1}{2}\mu_1^T\Sigma^{-1}\mu_1 + \ln\left(\frac{P(C_0)}{P(C_1)}\right)}_{\text{Constant } C}$$

transpose $x^T\Sigma^{-1}(\mu_0-\mu_1)$ as a number transpose is still a number

$= (\mu_0-\mu_1)^T(\Sigma^{-1})^T x + C$ * Σ is symmetric square matrix, $(\Sigma^{-1})^T = \Sigma^{-1}$.

$= (\mu_0-\mu_1)^T\Sigma^{-1}x + C$ Inverse of symmetric matrix, is still symmetric.

$w^T = (\mu_0-\mu_1)^T\Sigma^{-1}$

$b = -\frac{1}{2}\mu_0^T\Sigma^{-1}\mu_0 + \frac{1}{2}\mu_1^T\Sigma^{-1}\mu_1 + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$

$w = \Sigma^{-1}(\mu_0-\mu_1)$

$= \Sigma^{-1}\mu_0 - \Sigma^{-1}\mu_1$

(c) In (b), we modeled the class conditional distribution with same covariance matrix Σ. Now let us consider two classes that have difference covariance matrix as follows

$$P(x|C_0) = \frac{1}{(2\pi)^{n/2}|\Sigma_0|^{1/2}}\exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma_0^{-1}(x-\mu_0)\right),$$

$$P(x|C_1) = \frac{1}{(2\pi)^{n/2}|\Sigma_1|^{1/2}}\exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma_1^{-1}(x-\mu_1)\right).$$

Suppose we are able to find the maximum likelihood estimation of $\mu_0, \mu_1, \Sigma_0, \Sigma_1, P(C_0)$, and $P(C_1)$. Show that $a = x^TAx + w^Tx + b$ for some A, w and b. Find w and b in terms of $\mu_0, \mu_1, \Sigma_0, \Sigma_1, P(C_0)$, and $P(C_1)$. This shows that the decision boundary is quadratic.

$$a = \ln\left(\frac{\frac{1}{(2\pi)^{n/2}|\Sigma_0|^{1/2}}\cdot\exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma_0^{-1}(x-\mu_0)\right)\cdot P(C_0)}{\frac{1}{(2\pi)^{n/2}|\Sigma_1|^{1/2}}\cdot\exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma_1^{-1}(x-\mu_1)\right)\cdot P(C_1)}\right)$$

$$= \ln\left(\frac{1}{|\Sigma_0|^{1/2}}\cdot\exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma_0^{-1}(x-\mu_0)\right)\right) - \ln\left(\frac{1}{|\Sigma_1|^{1/2}}\cdot\exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma_1^{-1}(x-\mu_1)\right)\right) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= \ln\left(\frac{1}{|\Sigma_0|^{1/2}}\right) - \frac{1}{2}(x-\mu_0)^T\Sigma_0^{-1}(x-\mu_0) - \ln\left(\frac{1}{|\Sigma_1|^{1/2}}\right) + \frac{1}{2}(x-\mu_1)^T\Sigma_1^{-1}(x-\mu_1) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= -\frac{1}{2}(x^T-\mu_0^T)\Sigma_0^{-1}(x-\mu_0) + \frac{1}{2}(x^T-\mu_1^T)\Sigma_1^{-1}(x-\mu_1) + \ln\left(\frac{1}{|\Sigma_0|^{1/2}}\right) - \ln\left(\frac{1}{|\Sigma_1|^{1/2}}\right) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$$

$$= -\frac{1}{2}(x^T\Sigma_0^{-1} - \mu_0^T\Sigma_0^{-1})(x-\mu_0) + \frac{1}{2}(x^T\Sigma_1^{-1} - \mu_1^T\Sigma_1^{-1})(x-\mu_1) + C$$

$$= -\frac{1}{2}(x^T\Sigma_0^{-1}x - x^T\Sigma_0^{-1}\mu_0 - \mu_0^T\Sigma_0^{-1}x + \mu_0^T\Sigma_0^{-1}\mu_0) + \frac{1}{2}(x^T\Sigma_1^{-1}x - x^T\Sigma_1^{-1}\mu_1 - \mu_1^T\Sigma_1^{-1}x + \mu_1^T\Sigma_1^{-1}\mu_1) + C$$

$$= -\frac{1}{2}(x^T\Sigma_0^{-1}x - 2x^T\Sigma_0^{-1}\mu_0 + \mu_0^T\Sigma_0^{-1}\mu_0) + \frac{1}{2}(x^T\Sigma_1^{-1}x - 2x^T\Sigma_1^{-1}\mu_1 + \mu_1^T\Sigma_1^{-1}\mu_1) + C$$

$$= \frac{1}{2}x^T(\Sigma_1^{-1}-\Sigma_0^{-1})x + x^T\Sigma_0^{-1}\mu_0 - x^T\Sigma_1^{-1}\mu_1 \underbrace{-\frac{1}{2}\mu_0^T\Sigma_0^{-1}\mu_0 + \frac{1}{2}\mu_1^T\Sigma_1^{-1}\mu_1 + C}_{\text{constants}}$$

we transpose $x^T\Sigma_0^{-1}\mu_0$ and $x^T\Sigma_1^{-1}\mu_1$ as they are just numbers.

$$= x^T\left(\frac{1}{2}\Sigma_1^{-1} - \frac{1}{2}\Sigma_0^{-1}\right)x + \mu_0^T(\Sigma_0^{-1})^Tx - \mu_1^T(\Sigma_1^{-1})^Tx + C_1$$

$$= x^T\left(\frac{1}{2}\Sigma_1^{-1} - \frac{1}{2}\Sigma_0^{-1}\right)x + \left(\mu_0^T\Sigma_0^{-1} - \mu_1^T\Sigma_1^{-1}\right)x + C_1$$

$A = \left(\frac{1}{2}\Sigma_1^{-1} - \frac{1}{2}\Sigma_0^{-1}\right)$ $\quad w^T = \mu_0^T\Sigma_0^{-1} - \mu_1^T\Sigma_1^{-1}$ $\quad w = \Sigma_0^{-1}\mu_0 - \Sigma_1^{-1}\mu_1$

$b = -\frac{1}{2}\mu_0^T\Sigma_0^{-1}\mu_0 + \frac{1}{2}\mu_1^T\Sigma_1^{-1}\mu_1 + \ln\left(\frac{|\Sigma_1|^{1/2}}{|\Sigma_0|^{1/2}}\right) + \ln\left(\frac{P(C_0)}{P(C_1)}\right)$

4. We are given a training set $\{(x^{(i)}, y^{(i)}); i = \{1, \cdots, m\}\}$, where $x^{(i)} \in R^n$ and $y^{(i)} \in \{0, 1\}$. We consider the Gaussian Discriminant Analysis (GDA) model, which models $P(x|y)$ using multivariate Gaussian. Writing out the model, we have:

$$P(y = 1) = \phi = 1 - P(y = 0)$$

$$P(x|y = 0) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_0)^T \Sigma^{-1}(x - \mu_0)\right)$$

$$P(x|y = 1) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1)\right)$$

The log-likelihood of the data is given by:

$$L(\phi, \mu_0, \mu_1, \Sigma) = \ln P(x^{(i)}, \cdots, x^{(m)}, y^{(i)}, \cdots, y^{(m)}) = \ln \prod_{i=1}^{m} P(x^{(i)}|y^{(i)})P(y^{(i)}).$$

In this exercise, we want to maximize $L(\phi, \mu_0, \mu_1, \Sigma)$ with respect to ϕ, μ_0. The maximization over Σ is left for discussion.

(a) Write down the explicit expression for $P(x^{(1)}, \cdots, x^{(m)}, y^{(1)}, \cdots, y^{(m)})$ and $L(\phi, \mu_0, \mu_1, \Sigma)$.

$$P(x^1, \cdots, x^m, y^1, \cdots, y^m) = \prod_{i=1}^{m} \left[p(x_i|y_i=0)p(y_i=0)\right]^{1-y_i} \cdot \left[p(x_i|y_i=1)p(y_i=1)\right]^{y_i}$$

$$= \prod_{i=1}^{m} \left[\frac{1-\phi}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0)\right)\right]^{1-y_i} \cdot \left[\frac{\phi}{(2\pi)^{n/2}|\Sigma|^{1/2}} \cdot \exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1)\right)\right]^{y_i}$$

$$\ell(\phi, \mu_0, \mu_1, \Sigma) = \ln p(x^1, \cdots, x^m, y^1, \cdots, y^m)$$

$$= \sum_{i=1}^{m} (1-y_i)\ln\left(p(x_i|y=0)p(y_i=0)\right) + y_i \ln\left(p(x_i|y_i=1)p(y_i=1)\right)$$

$$= \sum_{i=1}^{m} (1-y_i)\ln\left(\frac{1-\phi}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x_i-\mu_0)^T\Sigma^{-1}(x_i-\mu_0)\right)\right) + y_i \cdot \ln\left(\frac{\phi}{2\pi^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x_i-\mu_1)^T\Sigma^{-1}(x_i-\mu_1)\right)\right)$$

$$= \sum_{i=1}^{m} (1-y_i) \cdot \left[\ln\left(\frac{1-\phi}{(2\pi)^{n/2}|\Sigma|^{1/2}}\right) - \frac{1}{2}(x_i-\mu_0)^T\Sigma^{-1}(x_i-\mu_0)\right] + y_i\left[\ln\left(\frac{1-\phi}{(2\pi)^{n/2}|\Sigma|^{1/2}}\right) - \frac{1}{2}(x_i-\mu_1)^T\Sigma^{-1}(x_i-\mu_1)\right]$$

$$= \sum_{i=1}^{m} (1-y_i) \cdot \left[\ln(1-\phi) - \frac{n}{2}\ln(2\pi) - \frac{1}{2}\ln(|\Sigma|) - \frac{1}{2}(x_i-\mu_0)^T\Sigma^{-1}(x_i-\mu_0)\right]$$
$$+ y_i\left[\ln(\phi) - \frac{n}{2}\ln(2\pi) - \frac{1}{2}\ln(|\Sigma|) - \frac{1}{2}(x_i-\mu_1)^T\Sigma^{-1}(x_i-\mu_1)\right]$$

(b) Find the maximum likelihood estimate for ϕ. How do you know such ϕ is the "best" but not the "worst"? Hint: Show that the second derivative of $L(\phi, \mu_0, \mu_1, \Sigma)$ with respect to ϕ is negative.

$$L(\phi, \mu_0, \mu_1, \Sigma) = \sum_{i=1}^{m} (1-y_i)\ln(1-\phi) + y_i \ln(\phi) + \cdots$$

$$= \ln(1-\phi) \cdot \sum_{i=1}^{m}(1-y_i) + \ln(\phi) \cdot \sum_{i=1}^{m}(y_i) + \cdots$$

$$\frac{\partial L}{\partial \phi} = \frac{-1}{1-\phi} \cdot \sum_{i=1}^{m}(1-y_i) + \frac{1}{\phi} \cdot \sum_{i=1}^{m}(y_i) = 0$$

$$\frac{1}{\phi} \cdot \sum_{i=1}^{m} y_i = \frac{1}{1-\phi} \sum_{i=1}^{m}(1-y_i) \qquad \text{let } S_1 = \sum_{i=1}^{m} y_i \quad S_0 = \sum_{i=1}^{m}(1-y_i)$$

$$\frac{S_1}{\phi} = \frac{S_0}{1-\phi} \qquad S_1 - S_1\phi = S_0\phi \qquad \frac{S_1}{S_0+S_1} = \phi$$

$$\frac{\partial^2 \ell}{\partial \phi^2} = -(1-\phi)^{-2} \cdot S_0 - \phi^{-2} \cdot S_1 = \frac{-1}{(1-\phi)^2} \cdot S_0 - \frac{1}{\phi^2} \cdot S_1 \leq 0$$

Since second derivative is always 0 or negative, we know that original function is always concave down, so when 1st derivative equals 0, must be a max.

(c) Find the maximum likelihood estimate for μ_0. How do you know such μ_0 is the "best" but not the "worst"? Hint: Show that the Hessian Matrix of $L(\phi, \mu_0, \mu_1, \Sigma)$ with respect to μ_0 is negative definite. You may use the following: if A is positive definite, then A^{-1} is also positive definite. Also B is negative definite if $-B$ is positive definite.

$$\ell(\phi, \mu_0, \mu_1, \Sigma) = \sum_{i=1}^{m} (1-y_i)\left[-\frac{1}{2}(x_i-\mu_0)^T\Sigma^{-1}(x_i-\mu_0)\right] + \dots$$

$$= \sum_{i=1}^{m} -\frac{1}{2}(1-y_i)\left[(x_i^T-\mu_0^T)\Sigma^{-1}(x_i-\mu_0)\right] + \dots$$

$$= \sum_{i=1}^{m} -\frac{1}{2}(1-y_i)\left(x_i^T\Sigma^{-1}x_i - 2\mu_0^T\Sigma^{-1}x_i + \mu_0^T\Sigma^{-1}\mu_0\right) + \dots$$

$$= \sum_{i=1}^{m} (1-y_i)\left(\mu_0^T\Sigma^{-1}x_i - \frac{1}{2}\mu_0^T\Sigma^{-1}\mu_0\right) + \dots$$

* Rules for Matrix Calculus:
  $\nabla_w(w^Tb) = b$
  $\nabla_w(w^TAw) = (A+A^T)w$
* $(\Sigma^{-1})^T = \Sigma^{-1}$ because Σ is symmetric.

$$\nabla_{\mu_0}\ell = \sum_{i=1}^{m}\left\{(1-y_i)(\Sigma^{-1}x_i)\right\} - \sum_{i=1}^{m}\left\{(1-y_i)\tfrac{1}{2}(\Sigma^{-1}+(\Sigma^{-1})^T)\mu_0\right\}$$

$$= \sum_{i=1}^{m}\left\{(1-y_i)(\Sigma^{-1}x_i)\right\} - \sum_{i=1}^{m}\left\{(1-y_i)\Sigma^{-1}\mu_0\right\}$$

$$= \sum_{i=1}^{m}\left\{(1-y_i)(\Sigma^{-1}x_i - \Sigma^{-1}\mu_0)\right\} = 0$$

$$\sum_{i=1}^{m}\left\{(1-y_i)(\Sigma^{-1}x_i)\right\} = \sum_{i=1}^{m}\left\{(1-y_i)\Sigma^{-1}\mu_0\right\}$$

$$\frac{\sum_{i=1}^{m}\left\{(1-y_i)(\Sigma^{-1}x_i)\right\}}{\sum_{i=1}^{m}\left\{(1-y_i)\Sigma^{-1}\right\}} = \mu_0$$

$$\mu_0 = \frac{\sum_{i=1}^{m}(1-y_i)x_i}{\sum_{i=1}^{m}(1-y_i)} = \frac{\sum_{i=1}^{m}(1-y_i)x_i}{S_0}$$

$$H_{\mu_0} = \nabla^2_{\mu_0}\ell = \sum_{i=1}^{m}(1-y_i)(-\Sigma^{-1}) = -S_0\Sigma^{-1}$$

From class, we know that Σ is a positive definite matrix, so Σ^{-1} is also positive definite. H_{μ_0} is negative definite since $-H_{\mu_0}$ is a positive definite matrix.

Since H_{μ_0} is negative definite, we know μ_0 found by setting gradient to zero must be a max.

5. In this exercise, you will implement a binary classifier using the Gaussian Discriminant Analysis (GDA) model in MATLAB. The data is given in *data.csv*. The first two columns are the feature values and the last column contains the class labels.

   (a) Visualization. Plot the data from different classes in different colors. Is the data linearly separable?

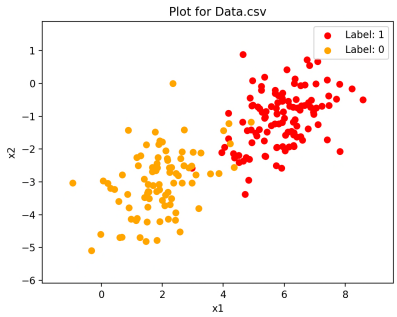

Data is not linearly Separable.

   (b) In the GDA model, we assume the class label follows a Bernoulli distribution and we model the class conditional distribution as multivariate Gaussian with same covariance matrix (Σ) and different means (μ_0 and μ_1). Find the maximum likelihood estimate of the parameters $P(y = 0)$ (parameter for the Bernoulli distribution), μ_0, μ_1 and Σ given this data set.

```
p(y = 0): 0.445
mu 0:
[[ 1.9195151]
 [-2.9972116]]
mu 1:
[[ 5.89818829]
 [-1.07926031]]
sigma:
[[1.0180746  0.38866005]
 [0.38866005 0.80363858]]
```

   (c) Using the result you find in Question 3 and your ML estimate of model parameters, find the decision boundary parameterized by $w^T x + b = 0$. Report w, b and plot the decision boundary on the same plot.

```
w:
[[-3.67554649]
 [-0.60899668]]
b: 12.90499335183926
```

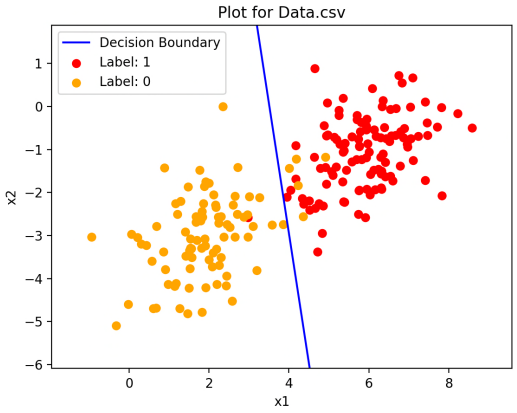

   (d) Visualize your results by plotting the contour of the two distributions $P(x, y = 0)$ and $P(x, y = 1)$. For consistency, set 'LevelList' ('level' for python) to logspace(-3,-1,7). Does your decision boundary pass through the points where the two distributions have equal probabilities ? Explain why.

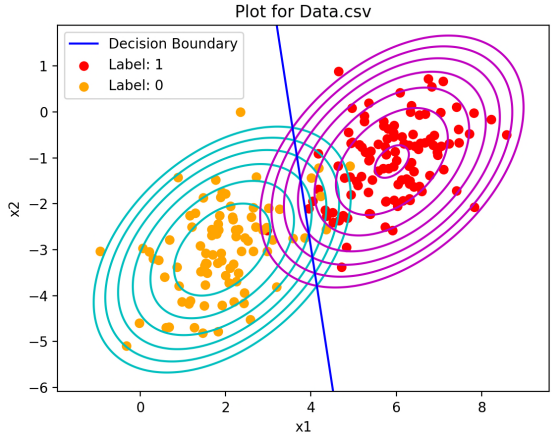

Yes, decision boundary passes through points where two distributions have equal probabilities because decision boundary is defined by

$$a = \ln\left(\frac{P(x|y_0)P(y_0)}{P(x|y_1)P(y_1)}\right) = \ln\left(\frac{P(x,y=0)}{P(x,y=1)}\right) = w^T x + b = 0$$

When $P(x,y=0) = P(x,y=1)$, is when two distributions have equal probability which means $a = \ln(1) = 0$, so whenever $P(x,y=0) = P(x,y=1)$ the points will be on the decision boundary.